

Tweet Sentiment Analysis

Dr. Jeffrey Strickland

9/13/2018

Sentiment Lexicons

One way to analyze the sentiment of a text is to consider the text as a combination of its individual words and the sentiment content of the whole text as the sum of the sentiment content of the individual words. This is often performed using lexicons. Linguistic theories generally regard human languages as consisting of two parts: a lexicon, essentially a catalogue of a language's words (its wordstock); and a grammar, a system of rules which allow for the combination of those words into meaningful sentences. Lexicons used of sentiment analysis contain sentiment words, like "trust" and "disgust," and all the words associated with a particular sentiment. These lexicons are "lined" together with the text under analysis by performing an inner-join with the text, the details of which we will discuss later.

Sentiment lexicons constructed via either crowdsourcing (using, for example, Amazon Mechanical Turk) or by the labor of one of the authors, and are validated using some combination of crowdsourcing again, restaurant or movie reviews, or Twitter data. Given this information, we may hesitate to apply these sentiment lexicons to styles of text dramatically different from what they were validated on, such as narrative fiction from 200 years ago. While it is true that using these sentiment lexicons with "Indian Philosophy," for example, may give us less accurate results than with tweets sent by a contemporary writer, we still can measure the sentiment content for words that are shared across the lexicon and the text.

The three general-purpose lexicons are - AFINN from Finn Årup Nielsen, - bing from Bing Liu and collaborators, and - nrc from Saif Mohammad and Peter Turney.

All three of these lexicons are based on unigrams, i.e., single words. These lexicons contain many English words and the words are assigned scores for positive/negative sentiment, and also possibly emotions like joy, anger, sadness, and so forth. The nrc lexicon categorizes words in a binary fashion ("yes"/"no") into categories of positive, negative, anger, anticipation, disgust, fear, joy, sadness, surprise, and trust. The bing lexicon categorizes words in a binary fashion into positive and negative categories. The AFINN lexicon assigns words with a score that runs between -5 and 5, with negative scores indicating negative sentiment and positive scores indicating positive sentiment. All of this information is tabulated in the sentiments dataset, and tidytext provides a function `get_sentiments()` to get specific sentiment lexicons without the columns that are not used in that lexicon.

Not every English word is in the lexicons because many English words are pretty neutral. It is important to keep in mind that these methods do not take into account qualifiers before a word, such as in “no good” or “not true”; a lexicon-based method like this is based on unigrams only.

Install Required Libraries

```
  if(!require(tidytext)) install.packages("tidytext")
## Loading required package: tidytext

  if(!require(tidyverse)) install.packages("tidyverse")
## Loading required package: tidyverse

## -- Attaching packages -----
----- tidyverse 1.2.1 -----

## v ggplot2 3.0.0      v purrr  0.2.5
## v tibble  1.4.2      v dplyr  0.7.6
## v tidyr   0.8.1      v stringr 1.3.1
## v readr   1.1.1      v forcats 0.3.0

## -- Conflicts -----
-- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

  if(!require(dplyr)) install.packages("dplyr")
  if(!require(tidyr)) install.packages("tidyr")
  if(!require(sentimentr)) install.packages("sentimentr")
## Loading required package: sentimentr

  if(!require(tm)) install.packages("tm")
## Loading required package: tm

## Loading required package: NLP

##
## Attaching package: 'NLP'

## The following object is masked from 'package:ggplot2':
##
##   annotate

  if(!require(readr)) install.packages("readr")
  if(!require(wordcloud)) install.packages("wordcloud")
## Loading required package: wordcloud

## Loading required package: RColorBrewer
```

```

    if(!require(lubridate)) install.packages("lubridate")
## Loading required package: lubridate
##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##     date

    if(!require(ggplot2)) install.packages("ggplot2")
    if(!require(ggraph)) install.packages("ggraph")
## Loading required package: ggraph

    if(!require(igraph)) install.packages("igraph")
## Loading required package: igraph
##
## Attaching package: 'igraph'
## The following objects are masked from 'package:lubridate':
##
##     %--%, union
## The following objects are masked from 'package:dplyr':
##
##     as_data_frame, groups, union
## The following objects are masked from 'package:purrr':
##
##     compose, simplify
## The following object is masked from 'package:tidyr':
##
##     crossing
## The following object is masked from 'package:tibble':
##
##     as_data_frame
## The following objects are masked from 'package:stats':
##
##     decompose, spectrum
## The following object is masked from 'package:base':
##
##     union

    if(!require(plotrix)) install.packages("plotrix")
## Loading required package: plotrix

```

Lexicon Example

The tidytext package contains several sentiment lexicons in the sentiments dataset.

```
library(tidytext)
sentiment

## function (text.var, polarity_dt = lexicon::hash_sentiment_jockers_rinker,
##      valence_shifters_dt = lexicon::hash_valence_shifters, hyphen = "",
##      amplifier.weight = 0.8, n.before = 5, n.after = 2, question.weight =
##      1,
##      adversative.weight = 0.85, neutral.nonverb.like = FALSE,
##      missing_value = 0, ...)
## {
##   UseMethod("sentiment")
## }
## <bytecode: 0x000000002113b470>
## <environment: namespace:sentiment>
```

We can look at the way specific lexicons score sentiment. The bing lexicon categorizes words in a binary fashion into positive and negative categories.

```
get_sentiments("bing")

## # A tibble: 6,788 x 2
##   word      sentiment
##   <chr>    <chr>
## 1 2-faced   negative
## 2 2-faces   negative
## 3 a+       positive
## 4 abnormal  negative
## 5 abolish  negative
## 6 abominable negative
## 7 abominably negative
## 8 abominate  negative
## 9 abomination negative
## 10 abort    negative
## # ... with 6,778 more rows
```

The AFINN lexicon assigns words with a score that runs between -5 and 5, with negative scores indicating negative sentiment and positive scores indicating positive sentiment.

```
get_sentiments("afinn")

## # A tibble: 2,476 x 2
##   word      score
##   <chr>    <int>
## 1 abandon     -2
## 2 abandoned   -2
## 3 abandons    -2
## 4 abducted    -2
```

```
## 5 abduction -2
## 6 abductions -2
## 7 abhor -3
## 8 abhorred -3
## 9 abhorrent -3
## 10 abhors -3
## # ... with 2,466 more rows
```

The nrc lexicon categorizes words in a binary fashion (“yes”/“no”) into categories of positive, negative, anger, anticipation, disgust, fear, joy, sadness, surprise, and trust.

```
get_sentiments("nrc")

## # A tibble: 13,901 x 2
##   word      sentiment
##   <chr>      <chr>
## 1 abacus    trust
## 2 abandon   fear
## 3 abandon   negative
## 4 abandon   sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,891 more rows
```

Example 1: The Inner Join

With data in a tidy format, sentiment analysis can be done as an inner-join. This is another of the great successes of viewing text mining as a tidy data analysis task; much as removing stop words is an anti-join operation, performing sentiment analysis is an inner-join operation. For this example, we will use a collection of tweets from President Donald Trump, as the lexicons are geared more toward the analysis of tweets than the analysis of ancient philosophies.

```
setwd("C:/Users/jeff/Documents/VIT_Course_Material/Data_Analytics_2018/data/")
trump_tweets<-read.csv("trump_tweets.csv",stringsAsFactor=FALSE)
str(trump_tweets)

## 'data.frame': 1050 obs. of 2 variables:
## $ time : chr "8/13/2018 8:57" "8/13/2018 9:21" "8/11/2018 1:28"
## $ tweets: chr " ....such wonderful and powerful things about me - a true
## Champion of Civil Rights - until she got fired. Omaro"| __truncated__ " While
## I know it's not presidential to take on a lowlife like Omarosa, and while I
## would rather not be doing "| __truncated__ " The big story that the Fake News
```

Media refuses to report is lowlife Christopher Steele's many meetings with De" | __truncated__ " The Republicans have now won 8 out of 9 House Seats, yet if you listen to the Fake News Media you would think " | __truncated__ ...

Observe that trump_tweets is already a data frame, and that there are 1050 tweets in the dataset, with line one as a header. We will look at the words with a "trust" score from the NRC lexicon. What are the most common "trust" words in Trump_Tweets? First, we need to take the text of the novels and convert the text to the tidy format using unnest_tokens().

The next code chunk converts the time to date in the dataset, which will be ordered by date.

```
trump_tweets$date<-as.Date(trump_tweets$time, "%m/%d/%Y %H:%M")
head(trump_tweets)

##           time
## 1 8/13/2018 8:57
## 2 8/13/2018 9:21
## 3 8/11/2018 1:28
## 4 8/8/2018 10:14
## 5  8/5/2018 7:49
## 6  8/5/2018 7:35
##
tweets
## 1 ....such wonderful and powerful things about me - a true Champion of
Civil Rights - until she got fired. Omarosa had Zero credibility with the
Media (they didn't want interviews) when she worked in the White House. Now
that she says bad about me, they will talk to her. Fake News! [Twitter for
iPhone] link
## 2      While I know it's not presidential to take on a lowlife like
Omarosa, and while I would rather not be doing so, this is a modern day form
of communication and I know the Fake News Media will be working overtime to
make even Wacky Omarosa look legitimate as possible. Sorry! [Twitter for
iPhone] link
## 3 The big story that the Fake News Media refuses to report is lowlife
Christopher Steele's many meetings with Deputy A.G. Bruce Ohr and his
beautiful wife, Nelly. It was Fusion GPS that hired Steele to write the phony
& discredited Dossier, paid for by Crooked Hillary & the DNC.... [Twitter for
iPhone] link
## 4                                The Republicans have now won 8 out
of 9 House Seats, yet if you listen to the Fake News Media you would think we
are being clobbered. Why can't they play it straight, so unfair to the
Republican Party and in particular, your favorite President! [Twitter for
iPhone] link
## 5
Too bad a large portion of the Media refuses to report the lies and
corruption having to do with the Rigged Witch Hunt - but that is why we call
them FAKE NEWS! [Twitter for iPhone] link
## 6 Fake News reporting, a complete fabrication, that I am concerned
about the meeting my wonderful son, Donald, had in Trump Tower. This was a
meeting to get information on an opponent, totally legal and done all the
```

time in politics - and it went nowhere. I did not know about it! [Twitter for iPhone] link

```
##           date
## 1 2018-08-13
## 2 2018-08-13
## 3 2018-08-11
## 4 2018-08-08
## 5 2018-08-05
## 6 2018-08-05
```

Next, we will tokenize the text.

```
trump_tweets<-trump_tweets %>% group_by(date) %>% mutate(ln=row_number())%>%
unnest_tokens(word,tweets) %>% ungroup()
head(trump_tweets,5)
```

```
## # A tibble: 5 x 4
##   time          date      ln word
##   <chr>        <date>  <int> <chr>
## 1 8/13/2018 8:57 2018-08-13     1 such
## 2 8/13/2018 8:57 2018-08-13     1 wonderful
## 3 8/13/2018 8:57 2018-08-13     1 and
## 4 8/13/2018 8:57 2018-08-13     1 powerful
## 5 8/13/2018 8:57 2018-08-13     1 things
```

Notice that we chose the name word for the output column from `unnest_tokens()`. This makes performing inner joins and anti-joins is thus easier because the sentiment lexicons and stop word datasets have columns named word.

Now that the text is in a tidy format with one word per row, we are ready to do the sentiment analysis. First, let's use the NRC lexicon and `filter()` for the "trust" words. Next, we will `filter()` the data frame with the text from the `trump_tweets` for the words and then use `inner_join()` to perform the sentiment analysis. What are the most common "trust" words in `trump_tweets`? We'll use `count()` from `dplyr` get this answer.

```
nrc_trust <- get_sentiments("nrc") %>%
  filter(sentiment == "trust")
trump_tweets %>%
  inner_join(nrc_trust) %>%
  count(word, sort = TRUE)
```

```
## Joining, by = "word"
```

```
## # A tibble: 189 x 2
##   word      n
##   <chr>  <int>
## 1 president 63
## 2 show      31
## 3 enjoy     25
## 4 good      20
## 5 vote      20
```

```
## 6 real 19
## 7 credibility 18
## 8 money 18
## 9 trade 17
## 10 white 17
## # ... with 179 more rows
```

We see mostly positive words associated with “trust”. Now we’ll look and see what “disgust” the President has.

```
nrc_disgust <- get_sentiments("nrc") %>%
  filter(sentiment == "disgust")
trump_tweets %>%
  inner_join(nrc_disgust) %>%
  count(word, sort = TRUE)

## Joining, by = "word"

## # A tibble: 126 x 2
##   word      n
##   <chr>   <int>
## 1 bad      66
## 2 dishonest 49
## 3 phony    25
## 4 witch    21
## 5 dying    19
## 6 collusion 18
## 7 enemy    17
## 8 john     16
## 9 terrible 12
## 10 hate    11
## # ... with 116 more rows
```

We can also examine how sentiment changes throughout the time period (10/19/2011 to 8/13/2018). We can do this with just a handful of lines that are mostly dplyr functions. First, we find a sentiment score for each word using the Bing lexicon and `inner_join()`. Next, we count up how many positive and negative words there are in defined in each tweet. We then use `spread()` so that we have negative and positive sentiment in separate columns, and lastly calculate a net sentiment (positive - negative).

```
library(tidyr)
trump_sentiment <- trump_tweets %>% inner_join(get_sentiments("bing"))

## Joining, by = "word"
```

Notice that we are plotting against the index on the x-axis that keeps track of narrative time in sections of text.








```
library(ggplot2)
trump_tweets%>%inner_join(get_sentiments("nrc")) %>%count(word,sentiment) %>%
group_by(sentiment)%>%top_n(10)%>%
```





```
    ungroup()
%>%mutate(word=reorder(word,n))%>%ggplot(aes(x=word,y=n,fill=sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ sentiment, scales = "free") +
  coord_flip()

## Joining, by = "word"

## Selecting by n
```




word



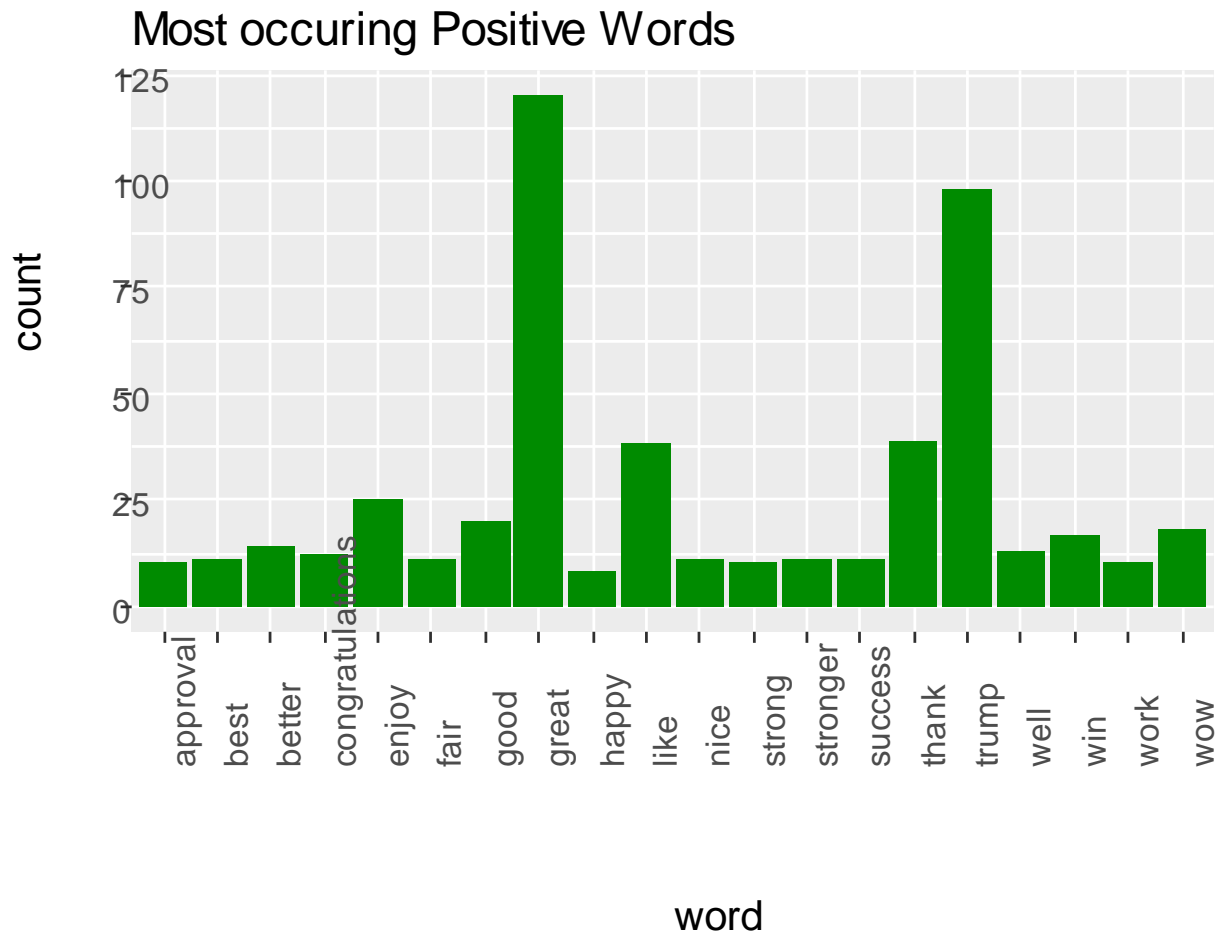
n

Positive & Negative Words over time

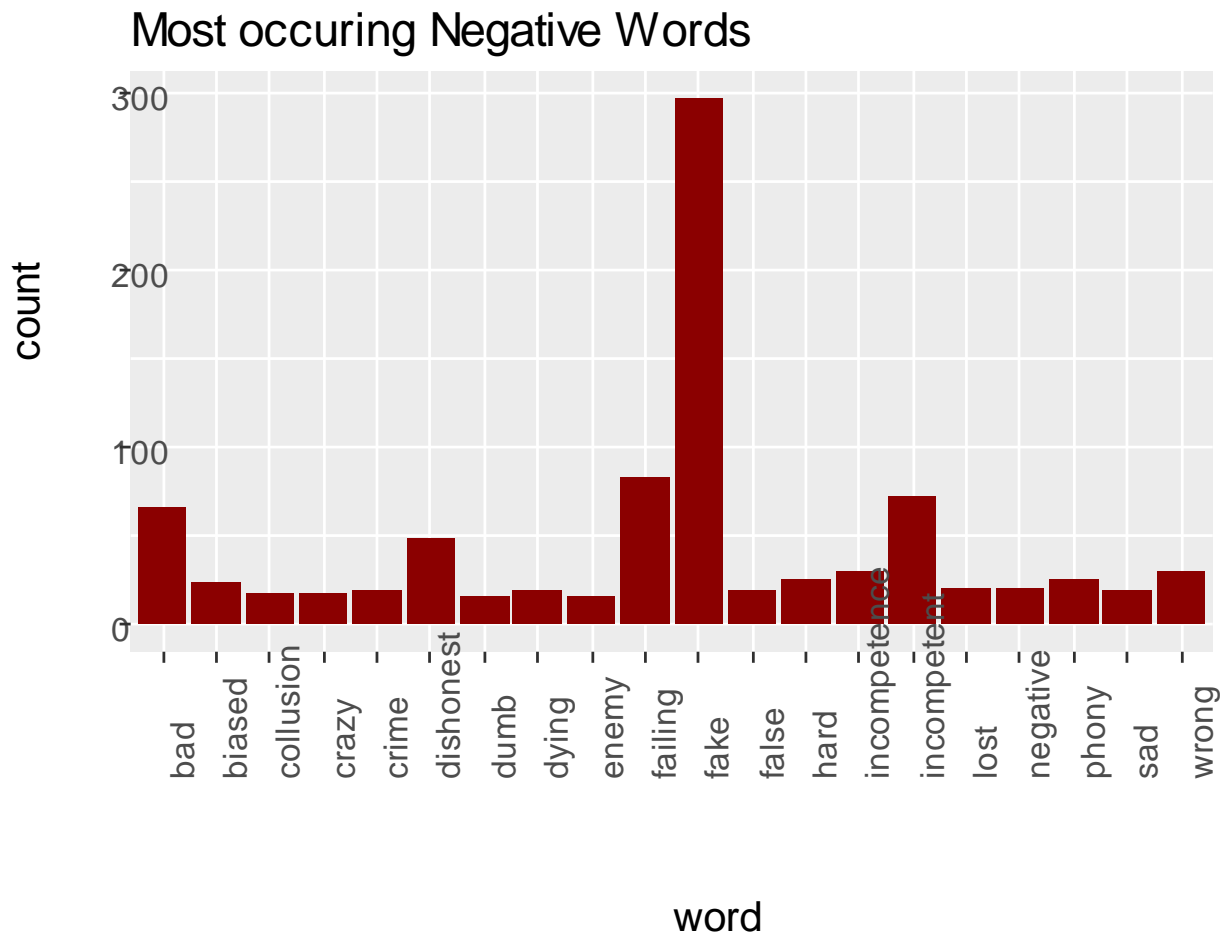
```
sentiment_by_time <- trump_tweets %>%  
  mutate(dt = floor_date(date, unit = "month")) %>%  
  group_by(dt) %>%  
  mutate(total_words = n()) %>%  
  ungroup() %>%  
  inner_join(get_sentiments("nrc"))  
  
## Joining, by = "word"  
  
sentiment_by_time %>%  
  filter(sentiment %in% c('positive','negative')) %>%  
  count(dt,sentiment,total_words) %>%  
  ungroup() %>%  
  mutate(percent = n / total_words) %>%  
  ggplot(aes(x=dt,y=percent,col=sentiment,group=sentiment)) +  
  geom_line(size = 1.5) +  
  geom_smooth(method = "lm", se = FALSE, lty = 2) +  
  expand_limits(y = 0)
```



```
pos_neg<-trump_sentiment %>%count(word,sentiment,sort=TRUE)
pos_neg %>% filter(sentiment=='positive')%>%head(20)
%>%ggplot(aes(x=word,y=n))+geom_bar(stat="identity",fill="green4")+
  theme(axis.text.x=element_text(angle=90))+labs(title="Most occuring
Positive Words",y="count")
```



```
pos_neg %>% filter(sentiment=='negative')%>%head(20)
%>%ggplot(aes(x=word,y=n))+geom_bar(stat="identity",fill="red4")+
  theme(axis.text.x=element_text(angle=90))+labs(title="Most occuring
Negative Words",y="count")
```



Most common positive and negative words

One advantage of having the data frame with both sentiment and word is that we can analyze word counts that contribute to each sentiment. By implementing `count()` here with arguments of both word and sentiment, we find out how much each word contributed to each sentiment.

```
bing_word_counts <- trump_tweets %>%
  inner_join(get_sentiments("bing")) %>%
  count(word, sentiment, sort = TRUE) %>%
  ungroup()
```

```
## Joining, by = "word"
```

```
bing_word_counts
```

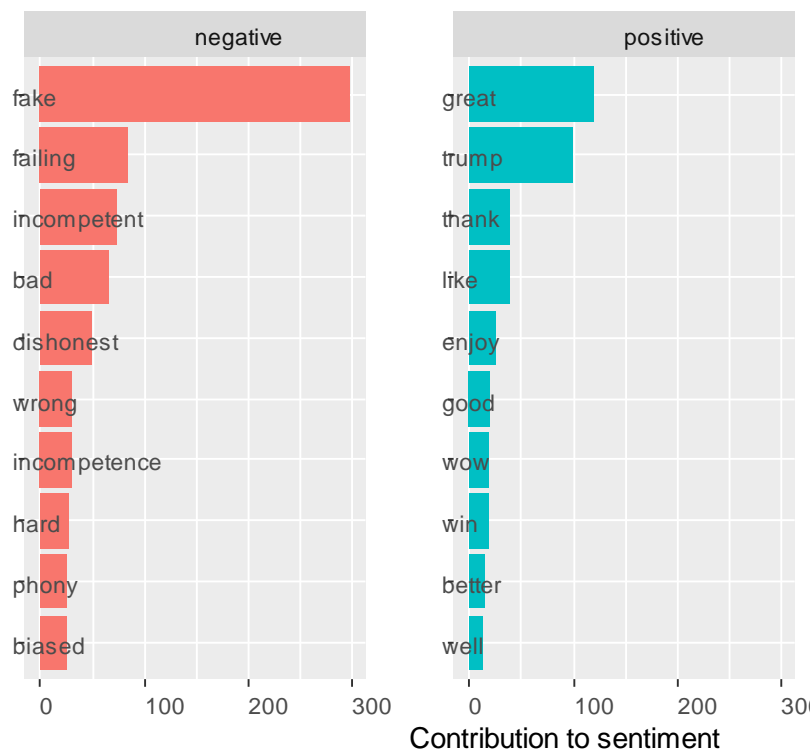
```
## # A tibble: 543 x 3
##   word      sentiment      n
##   <chr>      <chr>    <int>
## 1 fake      negative    298
## 2 great     positive    120
```

```
## 3 trump      positive  98
## 4 failing    negative  84
## 5 incompetent negative  73
## 6 bad        negative  66
## 7 dishonest  negative  49
## 8 thank      positive  39
## 9 like       positive  38
## 10 incompetence negative 31
## # ... with 533 more rows
```

This can be shown visually, and we can pipe straight into ggplot2, if we like, because of the way we are consistently using tools built for handling tidy data frames.

```
bing_word_counts %>%
  group_by(sentiment) %>%
  top_n(10) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(y = "Contribution to sentiment",
       x = NULL) +
  coord_flip()
```

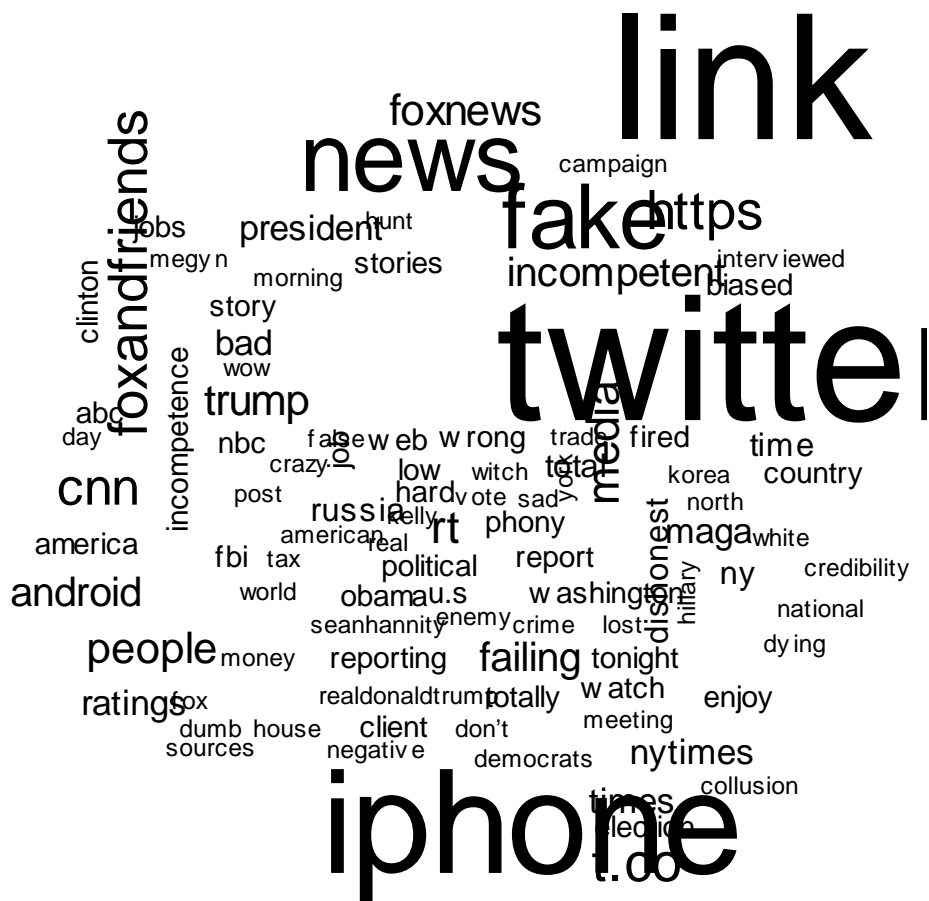
Selecting by n



Wordclouds

We've seen that this tidy text mining approach works well with `ggplot2`, but having our data in a tidy format is useful for other plots as well. For example, consider the `wordcloud` package, which uses base R graphics. Let's look at the most common words in `trump_tweets`, but this time as a wordcloud.

```
library(wordcloud)
trump_tweets %>%
  anti_join(stop_words) %>%
  count(word) %>%
  with(wordcloud(word, n, max.words = 100))
```



In other functions, such as `comparison.cloud()`, you may need to turn the data frame into a matrix with `reshape2`'s `acast()`. Let's do the sentiment analysis to tag positive and negative words using an inner join, then find the most common positive and negative words. Until the step where we need to send the data to `comparison.cloud()`, this can all be done with joins, piping, and `dplyr` because our data is in tidy format.

Word Frequencies

In this next section, we want to look at trump_tweets from the perspective of Mr. Trump's different roles, assuming that might be some sentiment differences between Candidate Trump and President Trump. The column "role" contains a flag to delineate Mr. Trump's different roles. After January 21, 2017, his role is "president" and prior to that, his role is "candidate."

The next code chunk removes user-made and standard English stopwords.

```
library(tm)
library(tidyr)
my_stops <- c("https", "a", "rt", "t.co", "for", "they", stopwords("en"))
trump_tweets <- trump_tweets %>%
  filter(!word %in% stop_words$word,
         !word %in% my_stops,
         !word %in% str_remove_all(stop_words$word, "''))
```

Now we can calculate word frequencies for each Donald Trump role. First, we group by person and count how many times each role used each word. Then we use left_join() to add a column of the total number of words used by each role.

```
frequency <- trump_tweets %>%
  group_by(role) %>%
  count(word, sort = TRUE) %>%
  left_join(trump_tweets %>%
            group_by(role) %>%
            summarise(total = n())) %>%
  mutate(freq = n/total)

## Joining, by = "role"

frequency

## # A tibble: 3,719 x 5
## # Groups:   role [2]
##   role      word      n total  freq
##   <chr>   <chr>   <int> <int> <dbl>
## 1 president twitter    584  9947 0.0587
## 2 president link      575  9947 0.0578
## 3 president iphone    541  9947 0.0544
## 4 president news     320  9947 0.0322
## 5 president fake     298  9947 0.0300
## 6 president foxandfriends 143  9947 0.0144
## 7 president media    109  9947 0.0110
## 8 president cnn      106  9947 0.0107
## 9 candidate link      91  3178 0.0286
## 10 president foxnews    89  9947 0.00895
## # ... with 3,709 more rows
```

This is a nice and tidy data frame but we would actually like to plot those frequencies on the x- and y-axes of a plot, so we will need to use `spread()` from `tidyr` to make a differently shaped data frame.

```
if(!require(tidyr)) install.packages("tidyr")
library(tidyr)
frequency <- frequency %>%
  select(role, word, freq) %>%
  spread(role, freq) %>%
  arrange(president, candidate)
frequency

## # A tibble: 3,210 x 3
##   word      candidate president
##   <chr>      <dbl>      <dbl>
## 1 18          0.000315  0.000101
## 2 admitted  0.000315  0.000101
## 3 agreed    0.000315  0.000101
## 4 allowing  0.000315  0.000101
## 5 articles  0.000315  0.000101
## 6 attempt   0.000315  0.000101
## 7 average   0.000315  0.000101
## 8 basic     0.000315  0.000101
## 9 bob       0.000315  0.000101
## 10 buy      0.000315  0.000101
## # ... with 3,200 more rows
```

Now this is ready for us to plot. We will use `geom_jitter()` so that we don't see the discreteness at the low end of frequency as much, and `check_overlap = TRUE` so the text labels don't all print out on top of each other (only some will print).

```
if(!require(scales)) install.packages("scales")

## Loading required package: scales

##
## Attaching package: 'scales'

## The following object is masked from 'package:plotrix':
##
##   rescale

## The following object is masked from 'package:purrr':
##
##   discard

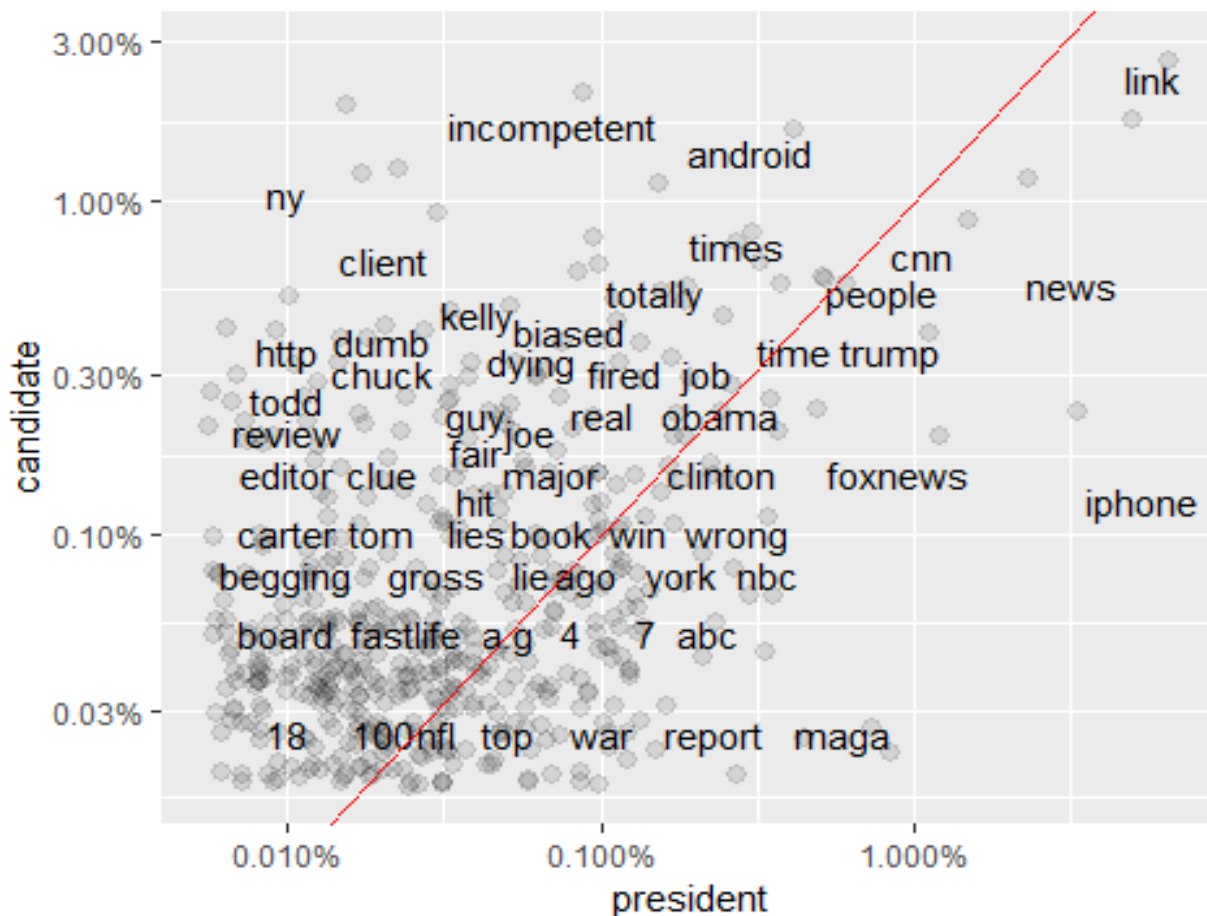
## The following object is masked from 'package:readr':
##
##   col_factor

library(scales)
ggplot(frequency, aes(president, candidate)) +
```

```
geom_jitter(alpha = 0.1, size = 2.5, width = 0.25, height = 0.25) +
geom_text(aes(label = word), check_overlap = TRUE, vjust = 1.5) +
scale_x_log10(labels = percent_format()) +
scale_y_log10(labels = percent_format()) +
geom_abline(color = "red")
```

```
## Warning: Removed 2701 rows containing missing values (geom_point).
```

```
## Warning: Removed 2701 rows containing missing values (geom_text).
```



Words near the line are used with about equal frequencies by President Trump and Candidate Trump, while words far away from the line are used much more by one person compared to the other. Words, hashtags, and usernames that appear in this plot are ones that we have both used at least once in tweets.

Comparing word usage

We just made a plot comparing raw word frequencies over our whole Twitter histories; now let's find which words are more or less likely to come from each Mr. Trump's roles, using the log odds ratio.

Here, we count how many times each role uses each word and keep only the words used more than 10 times. After a `spread()` operation, we can calculate the log odds ratio for each word, using

$\text{log odd ratio} = \ln(((n+1)/(\text{total}+1))_{\text{president}}/(((n+1)/(\text{total}+1))_{\text{candidate}}))$

where n is the number of times the word in question is used by each role and the total indicates the total words for each role.

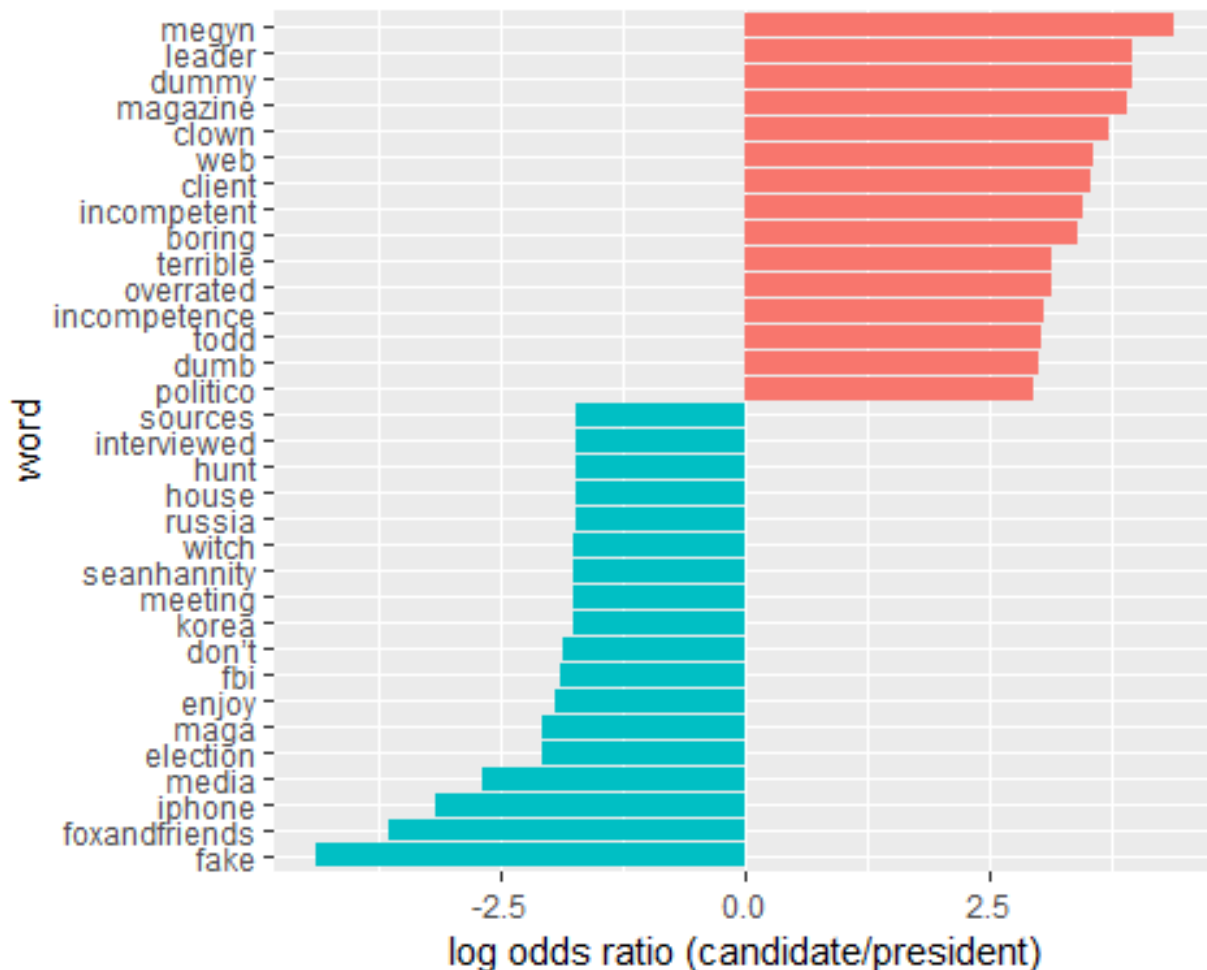
```
my_stops <- c("https", "http", "a", "rt", "t.co", "for", "they", "north",
"ny", stopwords("en"))
trump_tweets <- trump_tweets %>%
  filter(!word %in% stop_words$word,
         !word %in% my_stops,
         !word %in% str_remove_all(stop_words$word, ""))
word_ratios <- trump_tweets %>%
  filter(!str_detect(word, "^@")) %>%
  count(word, role) %>%
  group_by(word) %>%
  filter(sum(n) >= 10) %>%
  ungroup() %>%
  spread(role, n, fill = 0) %>%
  mutate_if(is.numeric, funs((. + 1) / (sum(.) + 1))) %>%
  mutate(logratio = log(candidate / president)) %>%
  arrange(desc(logratio))
word_ratios %>%
  arrange(abs(logratio))

## # A tibble: 203 x 4
##   word      candidate president logratio
##   <chr>      <dbl>      <dbl>    <dbl>
## 1 cnn         0.0191      0.0192 -0.00422
## 2 people      0.0150      0.0145  0.0330
## 3 polls       0.00205     0.00197  0.0371
## 4 china       0.00205     0.00215 -0.0499
## 5 coverage    0.00205     0.00215 -0.0499
## 6 million     0.00205     0.00215 -0.0499
## 7 campaign    0.00273     0.00287 -0.0499
## 8 world       0.00273     0.00287 -0.0499
## 9 clinton     0.00478     0.00449  0.0634
## 10 u.s        0.00478     0.00521 -0.0850
## # ... with 193 more rows
```

What are some words that have been about equally likely to come from Candidate Trump or President Trump?

```
word_ratios %>%
  group_by(logratio < 0) %>%
  top_n(15, abs(logratio)) %>%
  ungroup() %>%
  mutate(word = reorder(word, logratio)) %>%
```

```
ggplot(aes(word, logratio, fill = logratio < 0)) +
  geom_col(show.legend = FALSE) +
  coord_flip() +
  ylab("log odds ratio (candidate/president)") +
  scale_fill_discrete(name = "", labels = c("candidate", "president"))
```



Changes in word use

The section above looked at overall word use, but now let's ask a different question. Which words' frequencies have changed the fastest in Mr.Trumps Twitter feeds? Or to state this another way, which words has he tweeted about at a higher or lower rate as time has passed in his different roles? To do this, we will define a new time variable in the data frame that defines which unit of time each tweet was posted in. We can use `floor_date()` from `lubridate` to do this, with a unit of our choosing; using 1 month seems to work well for set of tweets.

After we have the time bins defined, we count how many times each of us used each word in each time bin. After that, we add columns to the data frame for the total number of words used in each time bin by each role and the total number of times each word was

used by each role We can then filter() to only keep words used at least some minimum number of times (30, in this case).

```
words_by_time <- trump_tweets %>%
  filter(!str_detect(word, "^@")) %>%
  mutate(time_floor = floor_date(date, unit = "1 month")) %>%
  count(time_floor, role, word) %>%
  group_by(role, time_floor) %>%
  mutate(time_total = sum(n)) %>%
  group_by(role, word) %>%
  mutate(word_total = sum(n)) %>%
  ungroup() %>%
  rename(count = n) %>%
  filter(word_total > 30)
words_by_time

## # A tibble: 482 x 6
##   time_floor role      word      count time_total word_total
##   <date>      <chr>    <chr>    <int>    <int>    <int>
## 1 2011-10-01 candidate incompetent    2        24        66
## 2 2011-10-01 candidate link          2        24        91
## 3 2011-11-01 candidate incompetent    1         9        66
## 4 2011-11-01 candidate link          1         9        91
## 5 2012-06-01 candidate link          1        20        91
## 6 2012-06-01 candidate twitter        1        20        86
## 7 2012-08-01 candidate incompetent    1        28        66
## 8 2012-08-01 candidate link          1        28        91
## 9 2012-09-01 candidate incompetent    4        49        66
## 10 2012-09-01 candidate link          4        49        91
## # ... with 472 more rows
```

Each row in this data frame corresponds to one role using one word in a given time bin. The count column tells us how many times that role used that word in that time bin, the time_total column tells us how many words that role used during that time bin, and the word_total column tells us how many times that person used that word over the whole year. This is the data set we can use for modeling.

We can use nest() from tidyr to make a data frame with a list column that contains little miniature data frames for each word. Let's do that now and take a look at the resulting structure.

```
nested_data <- words_by_time %>%
  nest(-word, -role)
nested_data

## # A tibble: 25 x 3
##   role      word      data
##   <chr>    <chr>    <list>
## 1 candidate incompetent <tibble [31 x 4]>
## 2 candidate link       <tibble [39 x 4]>
```

```
## 3 candidate twitter      <tibble [36 x 4]>
## 4 candidate android      <tibble [29 x 4]>
## 5 president cnn          <tibble [18 x 4]>
## 6 president failing      <tibble [15 x 4]>
## 7 president fake         <tibble [19 x 4]>
## 8 president foxnews      <tibble [19 x 4]>
## 9 president iphone       <tibble [19 x 4]>
## 10 president link        <tibble [20 x 4]>
## # ... with 15 more rows
```

This data frame has one row for each role-word combination; the data column is a list column that contains data frames, one for each combination of role and word. Let's use `map()` from `purrr` (Henry and Wickham 2018) to apply our modeling procedure to each of those little data frames inside our big data frame. This is count data so let's use `glm()` with `family = "binomial"` for modeling.

```
if(!require(purrr)) install.packages("purrr")
library(purrr)
nested_models <- nested_data %>%
  mutate(models = map(data, ~ glm(cbind(count, time_total) ~ time_floor, .,
                                family = "binomial")))
nested_models

## # A tibble: 25 x 4
##   role      word      data      models
##   <chr>    <chr>    <list>    <list>
## 1 candidate incompetent <tibble [31 x 4]> <S3: glm>
## 2 candidate link       <tibble [39 x 4]> <S3: glm>
## 3 candidate twitter    <tibble [36 x 4]> <S3: glm>
## 4 candidate android    <tibble [29 x 4]> <S3: glm>
## 5 president cnn        <tibble [18 x 4]> <S3: glm>
## 6 president failing    <tibble [15 x 4]> <S3: glm>
## 7 president fake       <tibble [19 x 4]> <S3: glm>
## 8 president foxnews    <tibble [19 x 4]> <S3: glm>
## 9 president iphone     <tibble [19 x 4]> <S3: glm>
## 10 president link      <tibble [20 x 4]> <S3: glm>
## # ... with 15 more rows
```

Now notice that we have a new column for the modeling results; it is another list column and contains `glm` objects. The next step is to use `map()` and `tidy()` from the `broom` package to pull out the slopes for each of these models and find the important ones. We are comparing many slopes here and some of them are not statistically significant, so let's apply an adjustment to the p-values for multiple comparisons.

```
if(!require(broom)) install.packages("broom")

## Loading required package: broom

library(broom)
library(tidyr)
slopes <- nested_models %>%
```

```
unnest(map(models, tidy)) %>%
filter(term == "time_floor") %>%
mutate(adjusted.p.value = p.adjust(p.value))
```

Now let's find the most important slopes. Which words have changed in frequency at a moderately significant level in our tweets?


```
top_slopes <- slopes %>%
  filter(adjusted.p.value < 0.05)
top_slopes
```

role <chr>	word <chr>	term <chr>	estimate <dbl>	std.error <dbl>
candidate	incompetent	time_floor	-0.0008457277	0.0002855362
candidate	link	time_floor	-0.0008922026	0.0002349048
candidate	twitter	time_floor	-0.0009304188	0.0002706097
candidate	android	time_floor	-0.0014821153	0.0004195812
president	failing	time_floor	-0.0042022246	0.0008896979
president	iphone	time_floor	-0.0010930502	0.0003032900
president	link	time_floor	-0.0013127521	0.0002774002
president	nytimes	time_floor	-0.0058712637	0.0010400612
president	twitter	time_floor	-0.0011893092	0.0002756080
president	foxandfriends	time_floor	-0.0034668078	0.0005628461

1-10 of 10 rows | 1-5 of 8 columns


To visualize our results, we can plot these words' use for both Candidate Trump and President Trump over tweet-period.

```
words_by_time %>%
  inner_join(top_slopes, by = c("word", "role")) %>%
  filter(role == "candidate") %>%
  ggplot(aes(time_floor, count/time_total, color = word)) +
  geom_line(size = 1.3) +
  labs(x = NULL, y = "Word frequency")
```

Now let's plot words that have changed frequency in President Trump's tweets.

```
words_by_time %>%
  inner_join(top_slopes, by = c("word", "role")) %>%
  filter(role == "president") %>%
  ggplot(aes(time_floor, count/time_total, color = word)) +
  geom_line(size = 1.3) +
  labs(x = NULL, y = "Word frequency")
```



Example 2: IsisFanboy

Sourcing Data

```
setwd("C:/Users/jeff/Documents/VIT_University/IsisFanboy/")
getwd()

## [1] "C:/Users/jeff/Documents/VIT_University/IsisFanboy"

tweets<-read.csv("AboutIsis.csv",stringsAsFactor=FALSE)
str(tweets)

## 'data.frame': 122619 obs. of 5 variables:
## $ name : chr "Sean Ferigan" "m.zakariyya" "ã\201jãððãððãððã\201ð"
## "chutney" ...
## $ username: chr "ferigan" "mzakariyya5" "yuzuchaaan777" "plainparatha"
## ...
## $ tweetid : num 7.52e+17 7.52e+17 7.52e+17 7.52e+17 7.52e+17 ...
## $ time : chr "7/11/2016 8:45:39 AM" "7/11/2016 8:45:39 AM" "7/11/2016
## 8:45:38 AM" "7/11/2016 8:45:38 AM" ...
## $ tweets : chr "ISIS influence on the decline as terrorists lose
## Twitter battles - CNET http://www.cnet.com/news/isis-influ|
## __truncated__ "RT @AyishaBaloch: #IndiaISISandBangladesh And world can ALSO
```

```
not ignore the truth revealing india 's role in pr"| __truncated__
"@Laika_isis @wink_BF
ãððãððãððãððãððã\201fã\201'|ã\201ªã\201ðã\201ðã\201ðã\201ð" "RT @KabirTaneja:
Anti-ISIS volunteer fighting with the Kurds. things are getting strange on
planet Earth. #Pok"| __truncated__ ...
```

Processing Data

View of data

```
tweets$date<-as.Date(tweets$time, "%d/%m/%Y %H:%M:%S")
head(tweets)
```

```
##          name          username          tweetid          time
## 1          Sean Ferigan          ferigan 7.524236e+17 7/11/2016 8:45:39 AM
## 2          m.zakariyya          mzakariyya5 7.524236e+17 7/11/2016 8:45:39 AM
## 3          ā\201jāāāāāāāā\201 yuzuchaaan777 7.524236e+17 7/11/2016 8:45:38
AM
## 4          chutney          plainparatha 7.524236e+17 7/11/2016 8:45:38 AM
## 5  à¥\220 àààààà°ààà à¥\220          dharam_vj 7.524236e+17 7/11/2016 8:45:37
AM
## 6 Dipendra Dipzo Khati DipendraDipzo 7.524236e+17 7/11/2016 8:45:36 AM
##
tweets
## 1
ISIS influence on the decline as terrorists lose Twitter battles          - CNET
http://www.cnet.com/news/isis-influence-twitter-on-the-decline-us-state-
department/#ftag=CAD590a51e
## 2
RT @AyishaBaloch: #IndiaISISandBangladesh And world can ALSO not ignore the
truth revealing india 's role in providin explosive to ISIS httpâ\200|
## 3
@Laika_isis @wink_BF
āāāāāāāāāāāāāāāāāāāā\201fā\201|ā\201ā\201ā\201ā\201ā\201ā\201ā
## 4
RT @KabirTaneja: Anti-ISIS volunteer fighting with the Kurds. things are
getting strange on planet Earth. #PokemonGO https://t.co/ARdBQ4â\200|
## 5 RT @MrsGandhi: It 's Urdu dailies not internet alone that 's turning
Muslims into terrorists #MustRead @tufaillelif
http://www.dailyo.in/politics/muslims-radicalisation-isis-hyderabad-ramzan-
internet-war-of-badr-prophet-muhammad-orlando-shooting/story/1/11599.html
## 6 RT @MrsGandhi: It 's Urdu dailies not internet alone that 's turning
Muslims into terrorists #MustRead @tufaillelif
http://www.dailyo.in/politics/muslims-radicalisation-isis-hyderabad-ramzan-
internet-war-of-badr-prophet-muhammad-orlando-shooting/story/1/11599.html
##          date
## 1 2016-11-07
## 2 2016-11-07
## 3 2016-11-07
## 4 2016-11-07
```

```
## 5 2016-11-07
## 6 2016-11-07
```

Tokenization

```
tweets$date<-as.Date(tweets$time, "%d/%m/%Y %H:%M:%S")
tidy_tweets<-tweets
%>%group_by(name,username,tweetid)%>%mutate(ln=row_number())%>%unnest_tokens(
word,tweets)%>%ungroup()
head(tidy_tweets,5)

## # A tibble: 5 x 7
##   name      username tweetid time          date      ln word
##   <chr>      <chr>    <dbl> <chr>          <date>    <int> <chr>
## 1 Sean Ferig~ ferigan  7.52e17 7/11/2016 8:45:3~ 2016-11-07      1 isis
## 2 Sean Ferig~ ferigan  7.52e17 7/11/2016 8:45:3~ 2016-11-07      1 influen~
## 3 Sean Ferig~ ferigan  7.52e17 7/11/2016 8:45:3~ 2016-11-07      1 on
## 4 Sean Ferig~ ferigan  7.52e17 7/11/2016 8:45:3~ 2016-11-07      1 the
## 5 Sean Ferig~ ferigan  7.52e17 7/11/2016 8:45:3~ 2016-11-07      1 decline

t_wrd<-tidy_tweets %>%count(word,sort=TRUE)
head(t_wrd,5)

## # A tibble: 5 x 2
##   word      n
##   <chr>  <int>
## 1 isis  113117
## 2 ã      93958
## 3 rt     86464
## 4 the    67765
## 5 à      58118
```

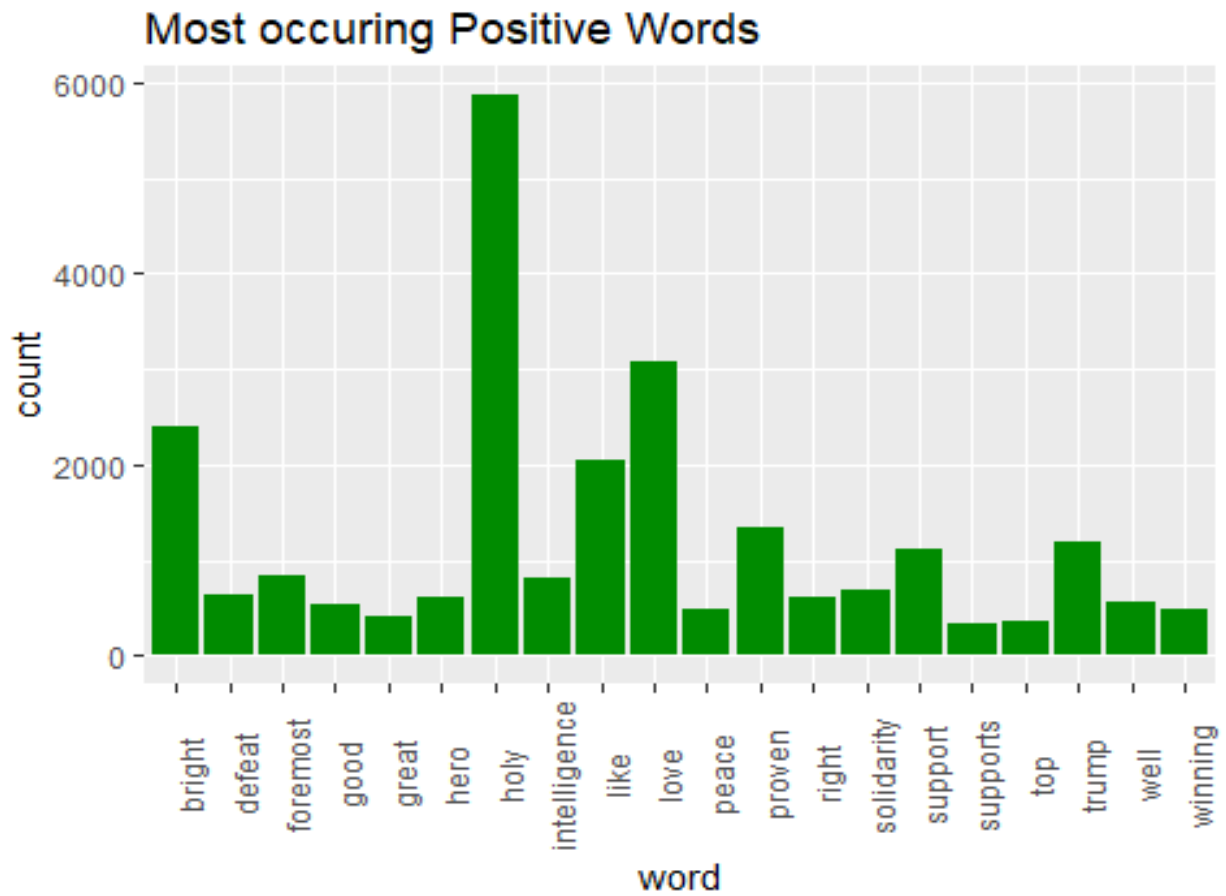
tidy_tweets contains all words like 'the','is','are' etc. So, we'll take only the sentimental words by joining with lexicons.

Combining with Lexicons to get Sentiment

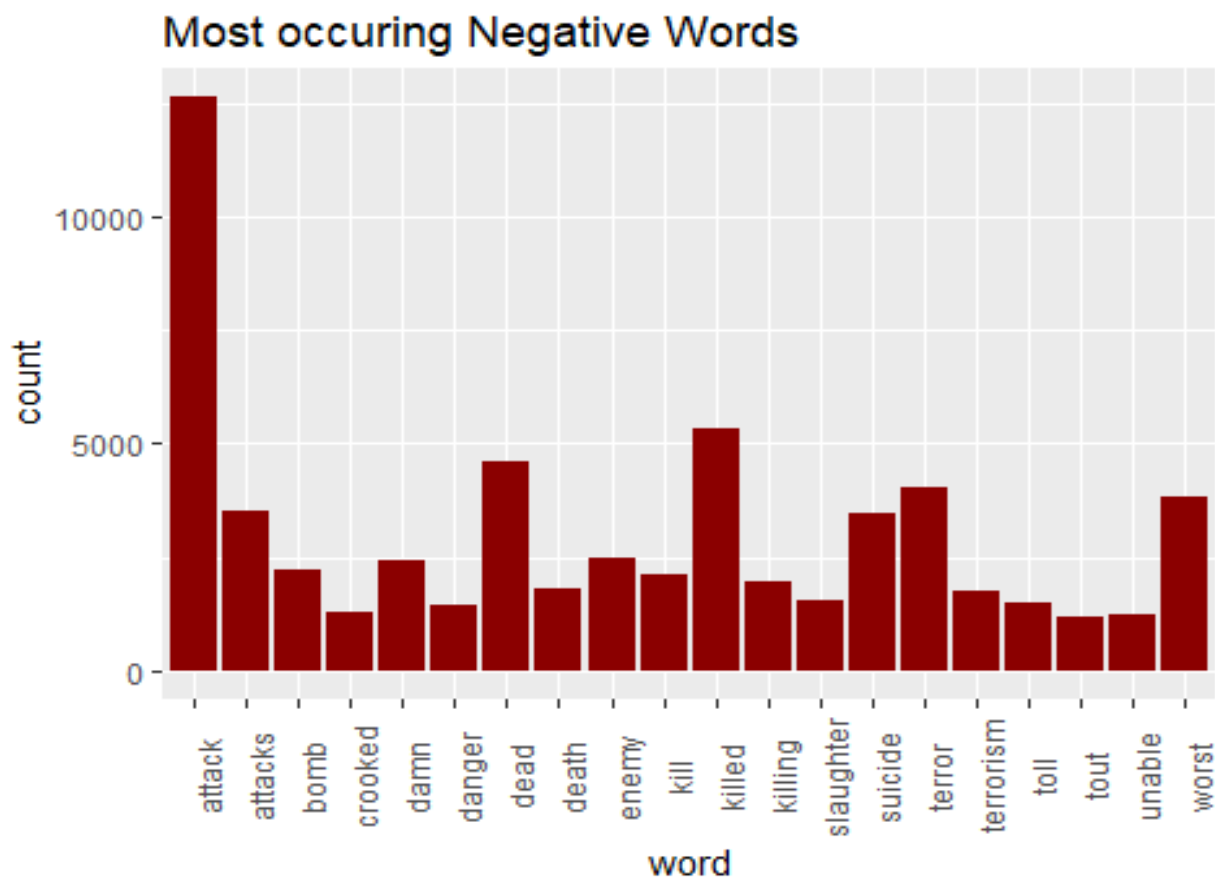
```
tweets_sentiment<-tidy_tweets%>% inner_join(get_sentiments("bing"))
## Joining, by = "word"
```

WordCloud of all words in ISIS tweets

```
tot_wrd<-tweets_sentiment%>%count(word,sort=TRUE)
wordcloud(tot_wrd$word,tot_wrd$n, min.freq =5, scale=c(5, .2), random.order =
FALSE, random.color = FALSE,colors = brewer.pal(8, "Dark2"))
```





```
pos_neg %>% filter(sentiment=='negative')%>%head(20)
%>%ggplot(aes(x=word,y=n))+geom_bar(stat="identity",fill="red4")+
  theme(axis.text.x=element_text(angle=90))+labs(title="Most occuring
Negative Words",y="count")
```




Percentage of positive and negative words in the tweets.

```
perc<-
tweets_sentiment%>%count(sentiment)%>%mutate(total=sum(n))%>%group_by(sentime
nt)%>%mutate(percent=round(n/total,2)*100)%>%ungroup()
label <-
  c( paste(perc$percent[1], '%', '-', perc$sentiment[1], sep=''),
    paste(perc$percent[2], '%', '-', perc$sentiment[2], sep=''))
pie3D(perc$percent, labels=label, labelcex=1.1, explode=0.1 )
```

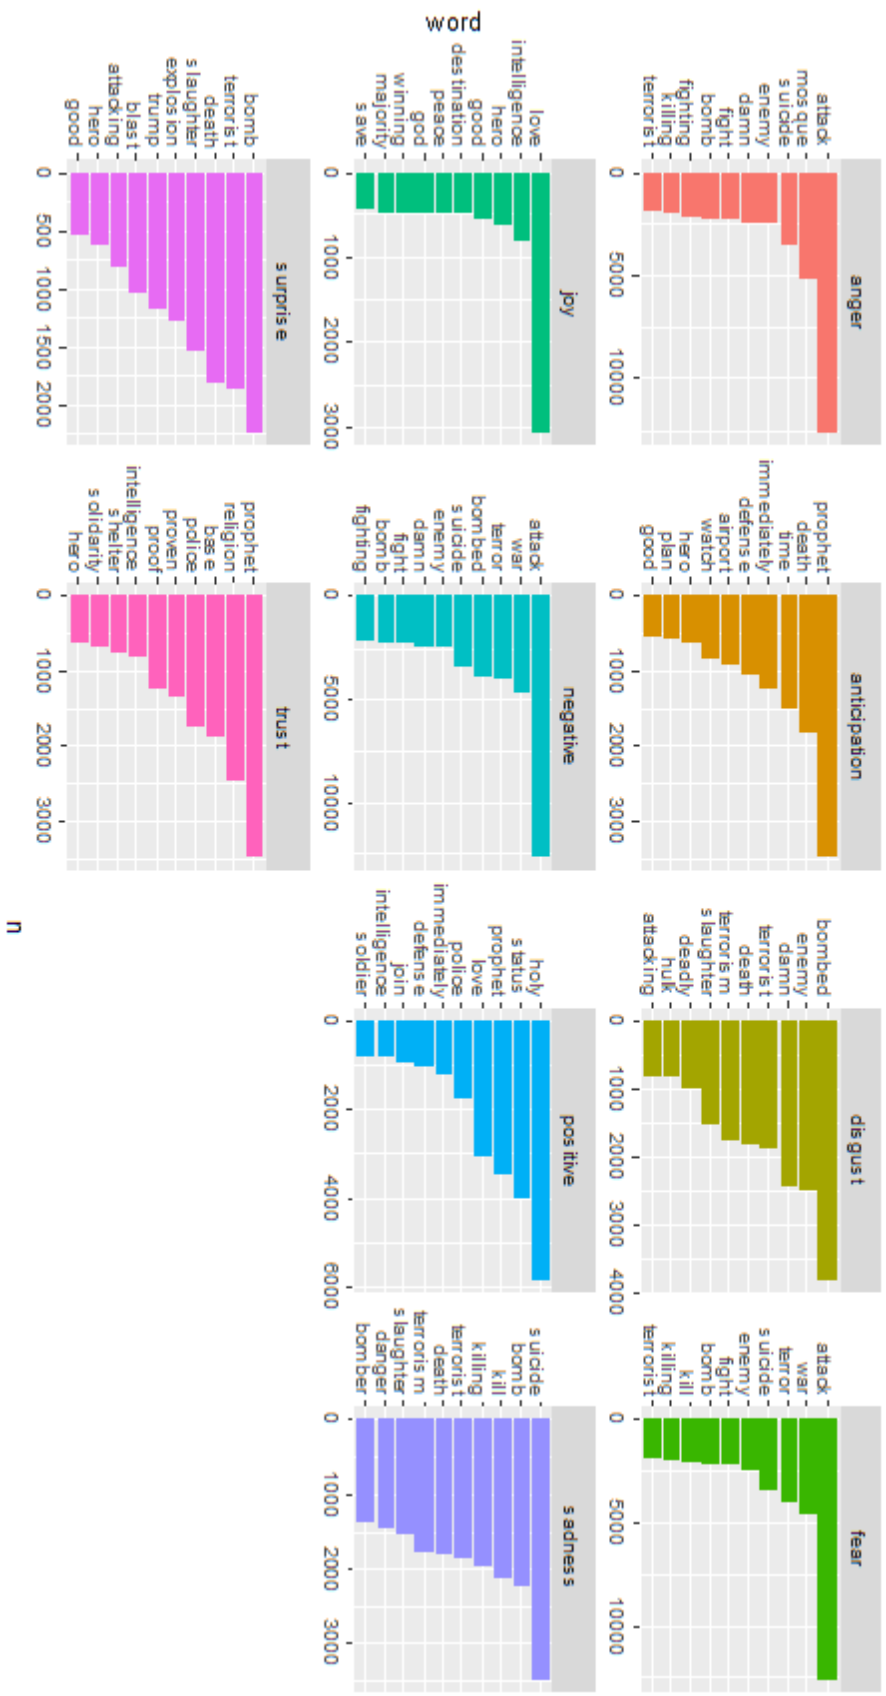


Wordcloud on Positive and Negative words

```
pos<-pos_neg %>% filter(sentiment=='positive')
neg<-pos_neg %>% filter(sentiment=='negative')
wordcloud(pos$word,pos$n, min.freq =5, scale=c(5, .2), random.order = FALSE,
random.color = FALSE,colors = brewer.pal(9,"Spectral"))
```




```
wordcloud(neg$word,neg$n, min.freq =5, scale=c(5, .2), random.order = FALSE,  
random.color = FALSE,colors = brewer.pal(9, "Spectral"))
```

n

Positive & Negative Words over Time

```
sentiment_by_time <- tidy_tweets %>%  
  mutate(dt = floor_date(date, unit = "month")) %>%  
  group_by(dt) %>%  
  mutate(total_words = n()) %>%  
  ungroup() %>%  
  inner_join(get_sentiments("nrc"))  
  
## Joining, by = "word"  
  
sentiment_by_time %>%  
  filter(sentiment %in% c('positive', 'negative')) %>%  
  count(dt, sentiment, total_words) %>%  
  ungroup() %>%  
  mutate(percent = n / total_words) %>%  
  ggplot(aes(x=dt, y=percent, col=sentiment, group=sentiment)) +  
  geom_line(size = 1.5) +  
  geom_smooth(method = "lm", se = FALSE, lty = 2) +  
  expand_limits(y = 0)
```



N-grams

An n-gram is a sequence of n “words” taken, in order, from a body of text. If we n = 2, then we get a bigram.

```
demo_bigrams <- unnest_tokens(tweets, input = tweets, output = bigram, token
= "ngrams", n=2)
demo_bigrams %>%
  count(bigram, sort = TRUE)

## # A tibble: 336,681 x 2
##   bigram          n
##   <chr>        <int>
## 1 à à          49079
## 2 ã ã          43140
## 3 islamic state 10743
## 4 2016 07       10569
## 5 isis is       9489
## 6 ø ø           8012
## 7 ã ã           7980
## 8 ø ù           7240
## 9 https t.co    6973
## 10 ù ø          6658
## # ... with 336,671 more rows

bigrams_separated <- demo_bigrams %>%
  separate(bigram, c("word1", "word2"), sep = " ")
bigrams_filtered <- bigrams_separated %>%
  filter(!word1 %in% stop_words$word) %>%
  filter(!word2 %in% stop_words$word)
```

New Bigram Counts

```
bigram_counts <- bigrams_filtered %>%
  count(word1, word2, sort = TRUE)
bigram_counts

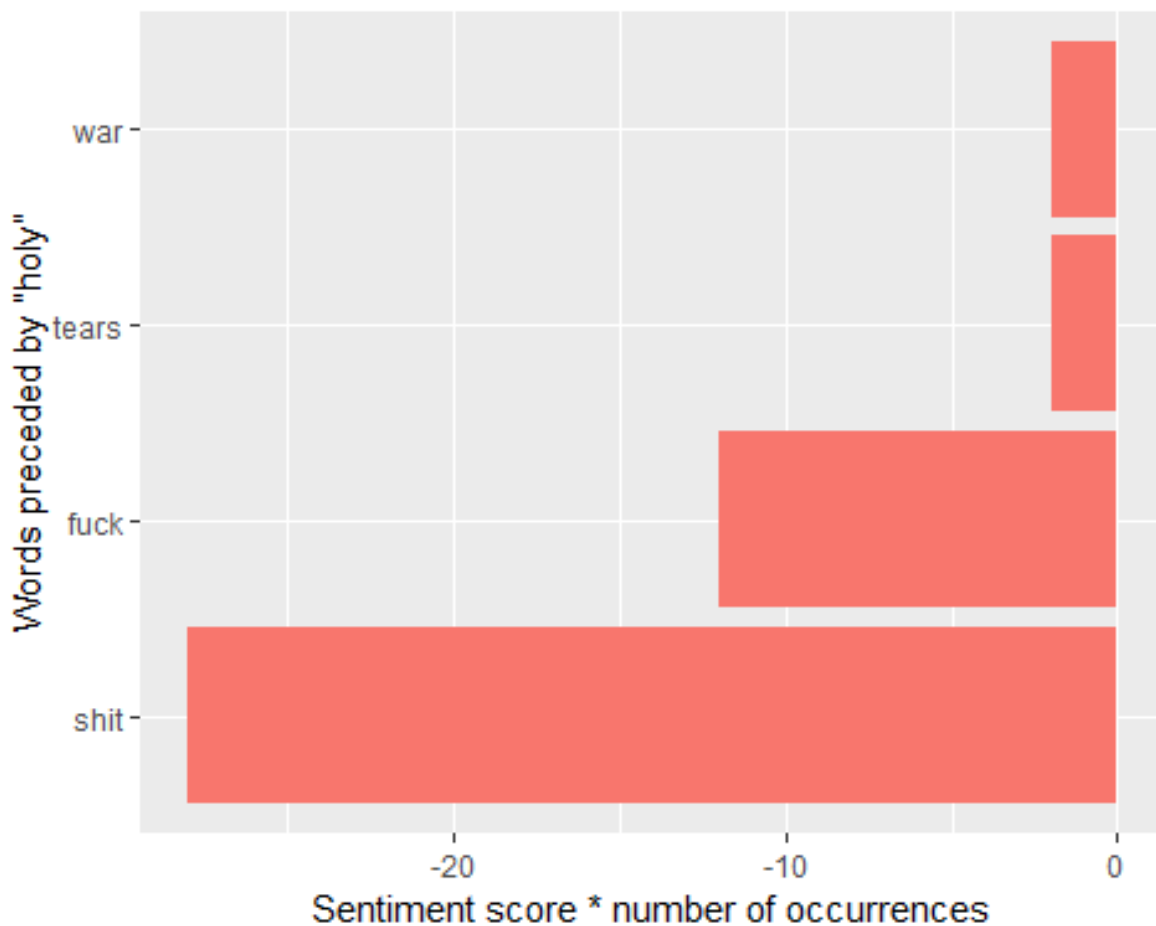
## # A tibble: 206,674 x 3
##   word1 word2      n
##   <chr> <chr> <int>
## 1 à     à     49079
## 2 ã     ã     43140
## 3 2016  07    10569
## 4 ø     ø     8012
## 5 ã     ã     7980
## 6 ø     ù     7240
## 7 https t.co   6973
## 8 ù     ø     6658
## 9 ã     ã     5950
## 10 å    ã     5617
## # ... with 206,664 more rows
```

Get AFINN Sentiments

```
AFINN <- get_sentiments("afinn")
not_words <- bigrams_separated %>%
  filter(word1 == "holy") %>%
  inner_join(AFINN, by = c(word2 = "word")) %>%
  count(word2, score, sort = TRUE) %>%
  ungroup()
```

Get Sentiment Scores

```
not_words %>%
  mutate(contribution = n * score) %>%
  arrange(desc(abs(contribution))) %>%
  head(20) %>%
  mutate(word2 = reorder(word2, contribution)) %>%
  ggplot(aes(word2, n * score, fill = n * score > 0)) +
  geom_col(show.legend = FALSE) +
  xlab("Words preceded by \"holy\"") +
  ylab("Sentiment score * number of occurrences") +
  coord_flip()
```



Plot Bigrams

```
bigram_graph <- bigram_counts %>%  
  filter(n > 10) %>%  
  graph_from_data_frame()  
  
## Warning in graph_from_data_frame(.): In `d` `NA` elements were replaced  
## with string "NA"  
  
bigram_graph  
  
## IGRAPH 147a33c DN-- 8600 17052 --  
## + attr: name (v/c), n (e/n)  
## + edges from 147a33c (vertex names):  
## [1] à ->à ã ->ã  
## [3] 2016 ->07 ø ->ø  
## [5] ã ->ã ù ->ù  
## [7] https ->t.co ù ->ø  
## [9] ã ->ã å ->å  
## [11] https ->twitter.com ù ->ù  
## [13] holy ->month ã ->ã  
## [15] ã ->å https ->å  
## + ... omitted several edges
```

Network of bigrams

```
set.seed(2017)  
ggraph(bigram_graph, layout = "fr") +  
  geom_edge_link() +  
  geom_node_point() +  
  geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```

